

Socket Programming

Unit 13 - Hands-On Networking - 2018

Prof. Dr.-Ing. Thorsten Herfet, [Andreas Schmidt](#), Pablo Gil Pereira

Telecommunications Lab, Saarland Informatics Campus, 23rd Feb. 2018

Recap

- IP addresses identify **hosts**.
- Ports identify **processes** on a host.
- TCP and UDP are **transport layer** protocols.
- **Services** to the user are provided by the **application layer**.
- **Client-Server** and **Peer-to-Peer** architectures can be used.

What is a Socket?

💡 **Intuition:** A `socket` is in the `kernel` one could `connect` with.

- Wall: Operating System
- Hole: API
- Do Stuff: Communicate Information (Send/Recv)
- Socket: OS Object



Source

Sockets are **the** network programming interface.

- Originally developed for UNIX, now everywhere.
- Well-defined network API for common network operations.
- Developing application layer protocols requires socket programming.

⚠️ **Lazy Programmers:** As TCP sockets are a bit complicated, devs tend to use HTTP for writing endpoints. It is easier, provides messages and there are frameworks.

Socket API

Abstractions

- API provides primitives (`bind`, `send`).
- Sockets abstract nasty details away:
 - Buffers incoming data.
 - Reliably transmit data.
 - Establish / teardown connections.
- Concepts shared amongst implementations.

Language Support

- Socket syntax differs between languages.
- We are using Python Low-Level Socket API.
- Python syntax similar to `glibc`.
 - Nice to use, as we can deal with objects in Python.
 - Don't have to mess around with file descriptors.

Learning Objectives

- How to use UDP, TCP, and raw sockets.

Why low-level API and no frameworks?

Dealing directly with socket API lets you understand...

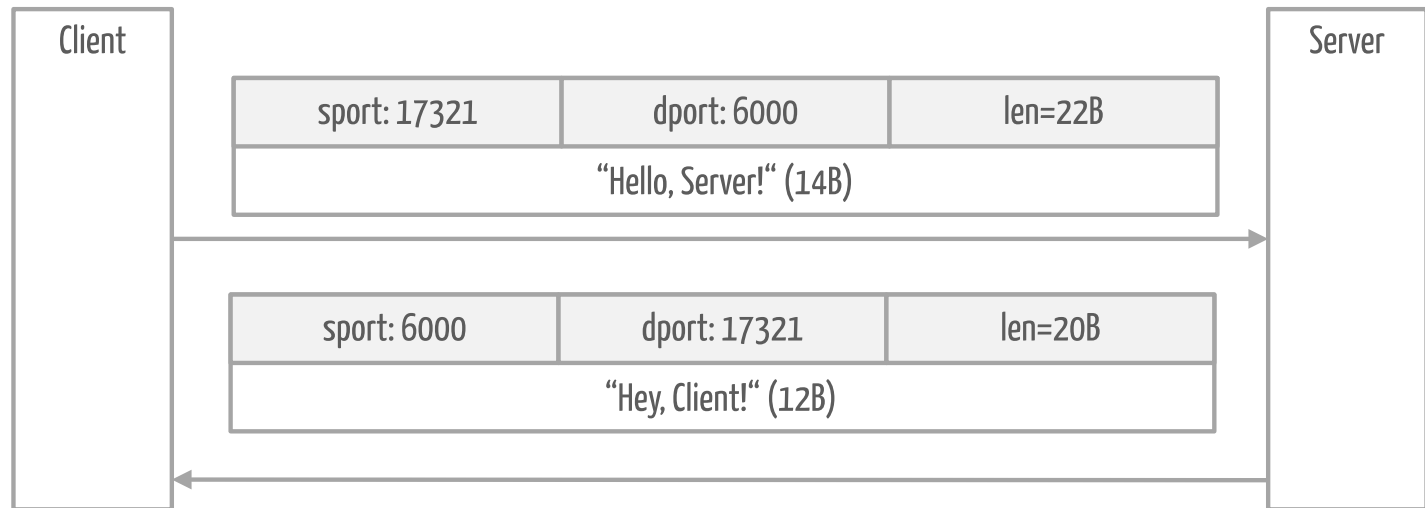
- ... the different states a socket can be in.
- ... how different parameters of sockets can be manipulated.
- ... what frameworks take over for you.
- ... what may be the reason when things don't work.

UDP Sockets

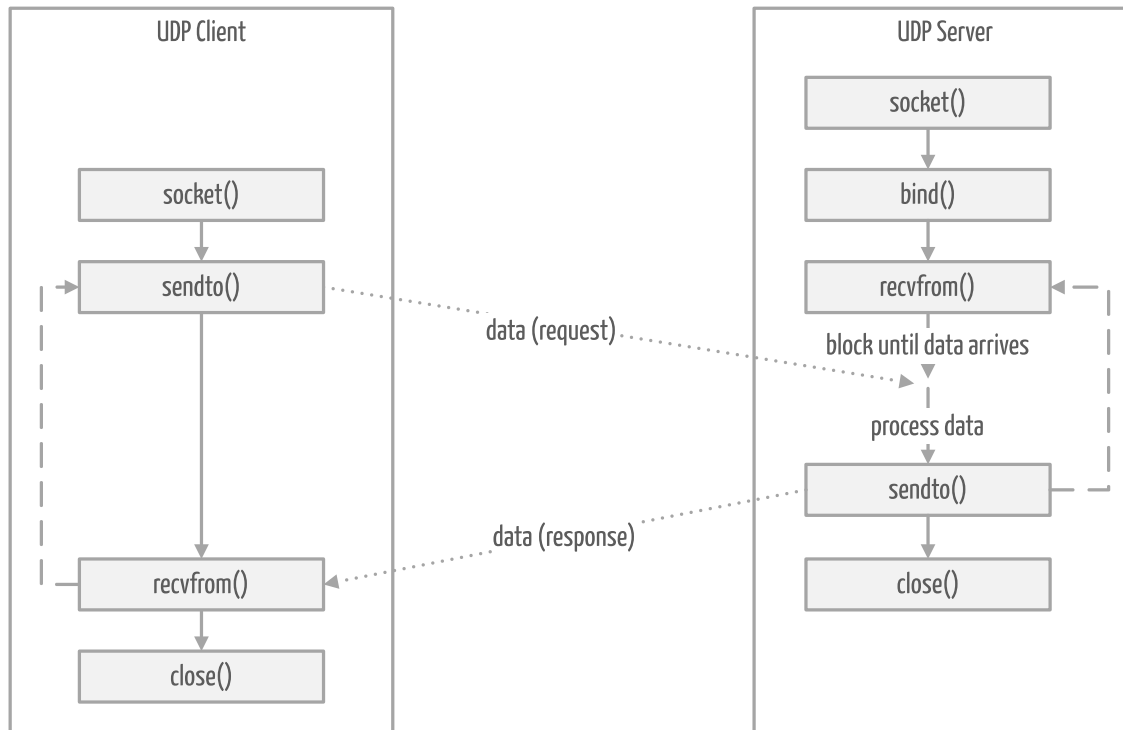
Motivating Example

© **Goal:** Simple UDP client/server programs.

- Client sends data (request).
- Server waits for request and responds (reply).



UDP | State Graph



UDP | Socket Creation (Step 1)

```
socket.socket(family, type, proto)
```

- Client and server create sockets.
- Returns: socket object
- **family**:
 - **AF_INET**: IPv4 socket
 - **AF_INET6**: IPv6 socket
 - **AF_PACKET**: device-level socket (receive raw packets)
 - **AF_UNIX**: UNIX socket
- **type**:
 - **SOCK_STREAM**: TCP socket
 - **SOCK_DGRAM**: UDP socket
 - **SOCK_RAW**: Raw socket
- **proto**: usually zero (0)

```
import socket
s = socket.socket(
    socket.AF_INET,      # family (IPv4)
    socket.SOCK_DGRAM,   # type (UDP)
    0)                   # proto (default)
```

Protocols

- The default of 0 lets the OS pick the appropriate **protocol**.
See IP protocol field in U10.
- Examples:
 - **IPPROTO_TCP**: 6
 - **IPPROTO_UDP**: 17
- All protocols can be found in **/etc/protocols** on UNIX.

UDP | Binding (Step 2)

```
s.bind(address_tuple)
```

- Server determines service port.
(mandatory step)
- Client can chose source port,
otherwise random
(optional step)
- `address_tuple`: Address and port:
 - IP address: specific IP, or
`0.0.0.0` to bind to all IPs.
 - Port: specific port, or 0 for
random port assignment.
- `bind()` can only be called once per
socket.
- Binding to ports ≤ 1023 requires
root privileges.

```
import socket
s = socket.socket(
    socket.AF_INET,      # family (IPv4)
    socket.SOCK_DGRAM,   # type (UDP)
    0)                  # proto (default)
s.bind(("192.168.0.1", 4567)) # (IP, port)
```

UDP | Receiving (Step 3a)

```
s.recvfrom(bufsize[, flags])
```

- Returns: (buffer, address) tuple
 - **buffer**: byte string with received data
 - **address**: address of sender (tuple of IP and port)
- **bufsize**: Maximum number of bytes to receive.
- **flags**: options (see later)
- Blocking, unless otherwise specified.
- Returns entire single UDP segment.
- Specify maximum UDP payload size as **bufsize** (round up to next 2^x)
- **bufsize** = 2^{16} = 65536
- or: pass **MSG_TRUNC** flag option

```
import socket
s = socket.socket(
    socket.AF_INET,          # family
    socket.SOCK_DGRAM,       # type
    0)                       # proto

s.bind(("192.168.0.1", 4567)) # (IP, port)

(data, address) = s.recvfrom(65536) # bufsize
print(f"Received {len(data)} bytes")
```

- Segment payload larger than **bufsize** is discarded.
- Example: **recvfrom(1024)** with payload of 2024 bytes will return first 1024B and discard remaining 1000B.

UDP | Send (Step 3b)

```
s.sendto(buf, flags, addr)
```

- Returns: number of bytes sent (`len(buf)`)
 - `sendto` atomic for UDP: entire buffer is sent (but no guarantee on reception)
- `buf`: byte string to be sent.
 - Has to fit into UDP payload of single IP packet.
 - Maximum size (IPv4): `65507`
- `flags`: optional control flags.
- `addr`: address tuple of recipient.


```
import socket
s = socket.socket(
    socket.AF_INET,          # family
    socket.SOCK_DGRAM,       # type
    0)                       # proto

s.bind(("192.168.0.1", 4567)) # (IP, port)

(data, address) = s.recvfrom(65536) # bufsize

data = data.decode("utf-8").toupper()

c = s.sendto(data.encode("utf-8"), # buf
             0,                   # flags
             address)              # addr
```

 **Multicast:** `addr` is specified when sending data and returned when receiving data, making UDP sockets usable for **one-to-many** transmissions.

UDP | Maximum Segment Size (IPv4)


UDP Length field is 16 bit large, hence maximum theoretical length is:

$$UDP_{max} = 2^{16} - 1 = 65535.$$

IP Total Length field is 16 bit large, hence maximum theoretical length is:

$$IP_{max} = 2^{16} - 1 = 65535 \text{ (too).}$$

UDP datagram must fit into one IP packet. Probably fragmented, but maximum length IP_{max} holds.

 Operating systems (Linux, etc.) know that you are putting UDP segments into IP packets.

IP ($IP_{hdr} \geq 20$) and UDP ($UDP_{hdr} = 8$) have a minimal / fixed header size.

Consequently, the practical maximum size of a UDP segment is:

$$IP_{max} - IP_{hdr} - UDP_{hdr} = 65535 - 20 - 8 = 65507$$

UDP | Close (Step 4)

```
s.close()
```

- Flushes all data if pending.
- Closing UDP sockets:
 - Free system resources.
 - Undo binding.
 - No explicit notification of remote side.

```
import socket
s = socket.socket(
    socket.AF_INET,          # family
    socket.SOCK_DGRAM,       # type
    0)                       # proto

s.bind(("192.168.0.1", 4567)) # (IP, port)

(data, address) = s.recvfrom(65536) # bufsize
data = data.decode("utf-8").toupper()

c = s.sendto(data.encode("utf-8"), # buf
             0,                   # flags
             address)             # addr

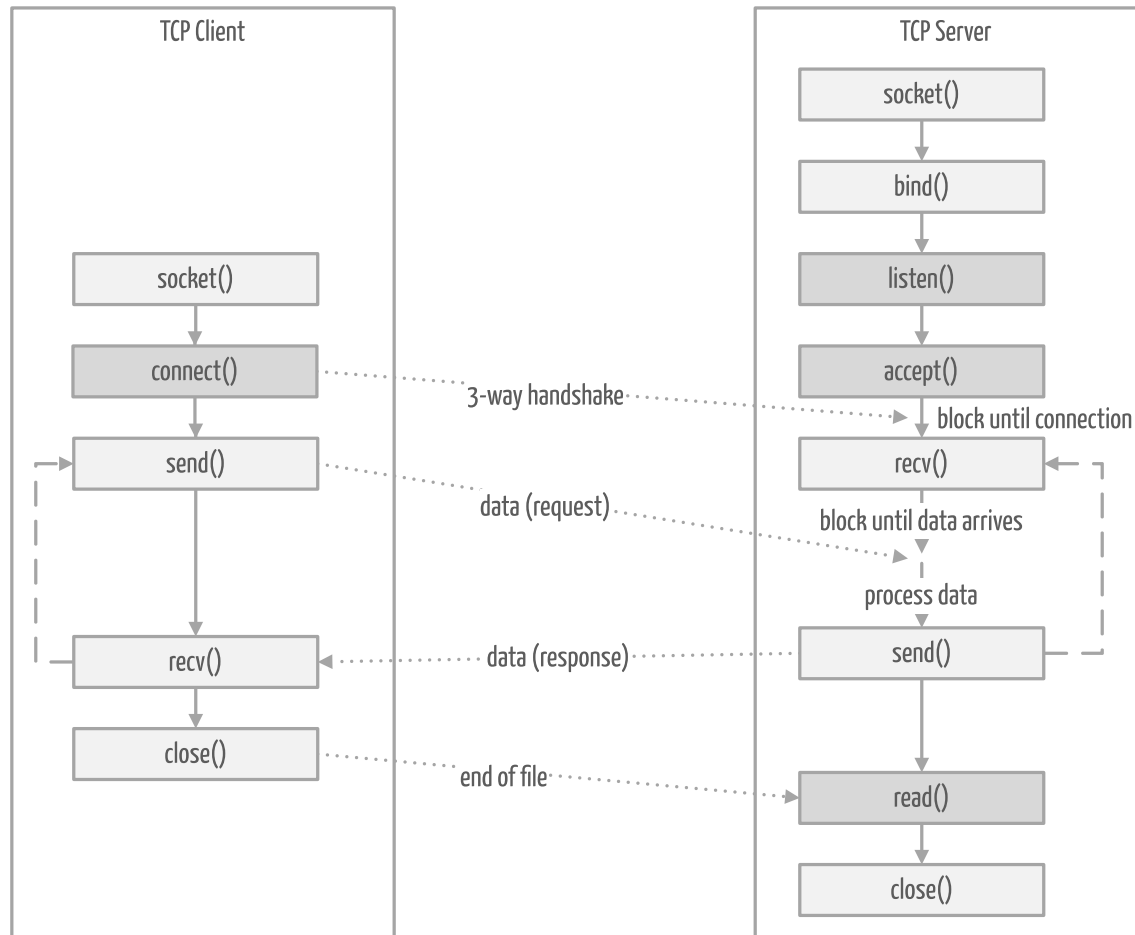
s.close()
```

TCP Sockets

TCP | Motivation

- TCP is
 - connection-oriented.
 - stream-oriented.
 - reliable.
- TCP sockets provide
 - similar functions to UDP.
 - additional connection management.
 - additionally message segmentation.

TCP | State Graph



TCP Server | Preparation (#1)

A) Create socket

- Command: `socket.socket(socket.AF_INET, socket.SOCK_STREAM, 0)`

B) Bind socket to port (similar to UDP)

- Command: `s.bind(address_tuple)`
- `address_tuple`: IP and port at which to wait for clients.

C) Listen for connections made to socket

- Starting at this command, kernel queues connections
- Command: `s.listen(backlog)`
- `backlog`: maximum number of queued connections

TCP Server | Connection Backlog

Concept

- After `listen`, kernel automatically completes TCP handshakes.
- OS kernel has two queues:
 - Incomplete connections: SYN received.
 - Completed connections: Handshake completed.
- `backlog`: maximum length of total entries (sum of both queues).
- Additional incoming TCP connections are ignored.

Design Decision

- Backlog of 0 is allowed, but can trigger unspecified behaviour.
- Backlog of 5 is a common value, but only suitable for small servers.
- Modern TCP-based servers have much larger values.

TCP Server | Accepting Connections (#2)

`s.accept()`

- Remove connection from completed `listen` queue.
 - First in, first out (FIFO) order.
- Returns: (conn, address) tuple.
 - `conn`: socket object of TCP connection.
 - `address`: address of sender (tuple of IP and port).
- `accept` is a blocking call:
 - Wait until client connects.
 - Blocks forever if nobody connects.

```
import socket
s = socket.socket(
    socket.AF_INET,      # family
    socket.SOCK_STREAM,  # type
    0)                   # proto

s.bind(("0.0.0.0", 4567)) # (IP, port)

s.listen(10)             # backlog

(conn, addr) = s.accept() # block until connect

(src_ip, src_port) = addr # connection accepted

print(f"Connection from {src_ip}:{src_port}")
```

Quiz

❓ Assume a socket has a backlog of 5 connections. What is the state after the following sequence of events?

- SYNs from 3 sources.
- Handshake completed with additional 2 sources.
- 1 socket `accept()`ed.

A: Backlog is full.

B: One slot free.

C: Two slots free.

D: Three slots free.

⚠️ Answer:

- ❌ A: Wrong, `accept()` removes one.
- ❌ C: Completed connections are not automatically removed.

- ✅ B: True, because 3 incompletes remain and one completed remains.
- ❌ D: Why? Off-by-two, maybe?

TCP Client | Establish Connection

A) Create Socket

- Command:
`socket.socket(socket.AF_INET,
socket.SOCK_STREAM, 0)`

```
import socket
s = socket.socket(
    socket.AF_INET,          # family
    socket.SOCK_STREAM,      # type
    0)                       # proto

s.connect(("127.0.0.1", 4567)) # (IP, port)

print("Successfully connected.")
```

B) Connect to server

- Command:
`s.connect(address_tuple)`
- `address_tuple`: IP and port of destination server
- Socket connected when `connect` returns.

TCP Sockets | Recv / Send (#3a/b)

Receiving

- Command:
`s.recv(bufsize[, flags])`
- Returns: Buffer that was received (up to `bufsize` bytes).
 - Empty buffer indicates other side **terminated** the connection.
- `bufsize`: Maximum number of bytes to receive.
- `flags`: options (see later)
- Blocking, unless otherwise specified.

Sending

- Command:
`s.send(buf[, flags])`
- Returns: Number of bytes sent.
 - `send` is **not** atomic for TCP: partial data might be sent.
- `buf`: Byte string to be sent.
 - Can be of any size, will be fragmented accordingly.
 - If too large, `send()` will only send **partial data**.
- `flags`: options (see later)

❓ For UDP, `recvfrom()` returned source address tuple. For TCP `recv()` does not. Why?

TCP Send/Recv | Correct Usage

- Issues with `recv` and `send`:
 - May return **less than** `bufsize` bytes.
 - `recv` may return multiple messages/packets.
 - `send` may not send all data passed as argument.
- Corner Cases:
 - You may need multiple `send` calls to transmit all data.
 - Multiple `send` calls may be handled by a single `recv` call.
 - A single `send` call may need multiple calls to `recv`.
- Cases may not be encountered under perfect conditions (e.g. testing), but it is quite common under high load.

Reliably *Send* all data

```
def send_all(s, buf):  
    buf_len = len(buf)  
    bytes_sent = 0  
    while bytes_sent < buf_len:  
        c = s.send(buf[bytes_sent:])  
        if c == 0:  
            raise RuntimeError("Socket " +  
                               "Connection Broken.")  
        bytes_sent += c
```

- `sendall()` in Python returns only if all data was sent (more convenient than C calls).

Reliably *Recv* all data

- Not trivial.
- Depends on message separation (e.g. `\n`, `\0` or `\n\r`).

TCP Send/Recv | Correct Usage II

💡 **Idea:** Reliably receive all data by using fixed message length.

Approach: Use `MSG_WAITALL`.

- Optional flag for `recv()`.
- `recv()` waits until `bufsize` bytes have been received.
- (... or a signal was caught.)
- (... or the connection was terminated)

⚡ **Problem:** Protocols with fixed length are usually bad.

❓ **Why is fixed length bad?**

A: Sending larger messages than predefined size is not easy.

B: Fixed size messages waste space, when the message is shorter than the fixed length.

C: Sender and recipient have to agree on length, but might have different opinions.

✅ A,B,C: True!

TCP Message Separation Options

Sender and receiver have to agree on message separation.

A) Fixed message length (e.g 1024 bytes).

- Easy to program, but lots of padding overhead.
- Large messages may not fit. **Bad design** in most cases.

B) Only one message per socket.

- Connect, send, recv, disconnect (e.g. basic HTTP used this).
- Connection establishment overhead.

C) Delimit messages by well-defined string/character (e.g. `\n`).

- Need escaping (or well-defined delimiter forbidden in message).
- Unclear how much needs to be read to reach delimiter.

D) Message length indicated at start of message.

- Acceptable overhead, very flexible.
- Message length field must have fixed size (therefore limits overall size).

TCP Sockets | Closing (#4)

```
s.close()
```

- Flushes all data if pending.
- Closing TCP sockets:
 - Send **FIN** segment ("I am done sending data.")
 - Other side calling **recv()** will get empty buffer.
 - At later point, remote will also close and send **FIN**.

Raw Sockets

Raw | Motivation

- TCP, UDP sockets:
 - Do most of the work for you (put data into a socket, give an address and the rest is figured out by the OS).
 - Required for transport layer networking.
- Sometimes, it is necessary to work with protocol on lowest layers:
 - ICMP: directly in IP without UDP/TCP.
 - ARP: Ethernet frame, without IP.
 - ...
- Security testing requires you to work on raw layer, so that packets can be crafted that would never be created in such a way by an OS.
- For more control and low-level networking: Use raw sockets!
 - Sender **can** generate arbitrary packets (including IP and Ethernet).
 - Sender **must** specify every detail of packets (including checksums).

Raw | Sending Packets

A) Create a socket with `socket.SOCK_RAW`.

- `socket.socket(socket.AF_PACKET, socket.SOCK_RAW, 0)`
- Note: Requires root/admin privileges (as we are circumventing the kernel/OS).

B) Bind to an interface (instead of IPv4 address).

- `s.bind(("eth0", 0))`

C) Generate and send packet.

- `s.send("A" * 20)` (malformed packet)
- Include Ethernet and IP header if needed.

Raw | Sending Packets (Convenient)

Generate packets using third-party library `dpkt`.

```
import dpkt

# create Ethernet, IP and UDP objects
eth = dpkt.ethernet.Ethernet(src="08:00:27:bf:17:b4", dst="ff:ff:ff:ff:ff:ff")
ip = dpkt.ip.IP(src="127.0.0.1", dst="127.6.6.6")
udp = dpkt.udp.UDP(sport=1234, dport=5678)

eth.data = ip          # set IP pkt as Ethernet payload
ip.data = udp          # set UDP segment to be IP pkt payload
ip.len = len(ip)       # set IP length
udp.ulen = len(udp)    # set UDP length

repr(ip)               # implicitly computes IP checksum
# Output:
# "IP(len=28, src='127.0.0.1', dst='127.6.6.6', data=UDP(sport=1234, dport=5678))"

repr(eth)              # show Ethernet packet
# Output:
#
# "Ethernet(dst='ff:ff:ff:ff:ff:ff', src='08:00:27:bf:17:b4',
#      data=IP(len=28, src='127.0.0.1', dst='127.6.6.6',
#      data=UDP(sport=1234, dport=5678)))"
```

Raw | Receiving Packets

🎯 **Goal:** Receive all packets on interface.

- Create socket with `socket.SOCK_RAW`:

```
ETH_P_ALL = socket.ntohs(0x0003) # all Ethernet protos
socket.socket(socket.AF_PACKET, # domain
              socket.SOCK_RAW, # type
              ETH_P_ALL)        # protocol (non-empty / non-default)
```

- Put NIC to **promiscuous mode** and even get packets not intended for you.
- Receive packet:
 - `s.recv(4096)`
 - Return Ethernet frame as captured on interface.
 - Note: captures packets even if they are filtered by firewall.

Raw | Receiving Packets (Convenient)

Parse packets using third-party library `dpkt`.

```
import dpkt

pkt = s.recv(4096)          # receive via raw socket

dpkt.ethernet.Ethernet(pkt)  # print dpkt parsing result

# Output
# Ethernet(dst="\x08\x00'\xbf\x17\xb4", src='RT\x00\x125\x02',
#          data=IP(len=40, id=17571, p=6, sum=7709, src='\n\x00\x02\x02', dst='\n\x00\x02\x0f', opts='',
#          data=TCP(sport=54945, dport=22, seq=126587142, ack=4007547311, flags=16, sum=59372, opts=''))))

eth = dpkt.ethernet.Ethernet(pkt)  # get Ethernet frame object
ip = eth.data                      # get IP packet object
tcp = ip.data                      # get TCP packet object
tcp.dport # access member variables
# Output:
# 22
```

Raw | Power

” “With Great Power Comes Great Responsibility” -

🔍 Network Link Sniffing

- See every packet on the wire.
And many of those in the air.
- Passive attacks such as local eavesdropping become easy.

📦 Arbitrary Packet Injection

- You can inject completely arbitrary packets.
- "Spoof" any arbitrary MAC or IP address(es).
- Low resource requirements (in contrast to full TCP/IP stack).

Wrap-Up

Questions?

Take-Home Messages

- **Sockets** are **the API** for network programming and an important concept.
- **UDP socket** offer atomic send/receive operations.
- **TCP sockets** require more attention (connection handling, partial receive/send).
- **Raw sockets** give your complete control, but also put burden on you.

Further Reading

- [Python 3.6 - Low-level networking interface](#)
- [Python 3.6 - Struct module](#)
- Stevens, Fenner, Rudoff: UNIX Network Programming, Volume 1
 - Book examples in C, we use Python.