

# Socket Debugging, Server Design and Special Functions

Unit 14 - Hands-On Networking - 2018

*Prof. Dr.-Ing. Thorsten Herfet, [Andreas Schmidt](#), Pablo Gil Pereira*

**Telecommunications Lab, Saarland Informatics Campus, 23rd Feb. 2018**

# Socket Debugging Tools

- **Validate binding socket works:** `netstat -tulpn`
  - `-t` to show TCP sockets.
  - `-u` to show UDP sockets.
  - `-n` to show numeric IP addresses.
  - `-l` to show listening sockets.
  - `-p` to also show the process.
- **Validate TCP server working:** `netcat <server> <port>`
  - Go ahead typing your data. Enter to send.
- **Validate UDP server working:** `netcat -u <server> <port>`
  - Go ahead typing your data. Enter to send.

# Return To Sender

⚡ Sometimes you might contact a server that is not listening.

## Reactions

**TCP:** (RST=1).

**UDP:** ICMP

.



# Servers

**(The Good, The Bad and The Ugly)**

# UDP Time Service

```
# Server

from socket import *
import datetime

s = socket(AF_INET, SOCK_DGRAM)
s.bind(("0.0.0.0", 9000))
req, sender = s.recvfrom(1000)
print("Sender", sender, "asked for:",
      req.decode("utf-8"))
now = datetime.datetime.now()
print("Replying with:", now)
s.sendto(now.isoformat().encode("utf-8"), sender)
s.close()
```

```
# Client

from socket import *

s = socket(AF_INET, SOCK_DGRAM)
cmd = "time"
c = s.sendto(cmd.encode("utf-8"),
             ('127.0.0.1', 9000))
print("Requested",cmd,"which is",c,"bytes long.")
time, sender = s.recvfrom(100)
print("Sender",sender,"replied with time:",
      time.decode("utf-8"))
s.close()
```

**❓ What is the problem with this server implementation?**

Single-use, disposable. Restart server after each served response.

# UDP Time Service - Iterative

```
# Server

from socket import *
import datetime

s = socket(AF_INET, SOCK_DGRAM)
s.bind(("0.0.0.0", 9000))

try:
    while True:
        req, sender = s.recvfrom(1000)
        print("Sender", sender, "asked for:",
              req.decode("utf-8"))
        now = datetime.datetime.now()
        print("Replying with:", now)
        s.sendto(now.isoformat().encode("utf-8"),
                  sender)
except KeyboardInterrupt:
    pass
finally:
    s.close()
```

```
# Client

from socket import *

s = socket(AF_INET, SOCK_DGRAM)
cmd = "time"
c = s.sendto(cmd.encode("utf-8"),
              ('127.0.0.1', 9000))
print("Requested",cmd,"which is",c,"bytes long.")
time, sender = s.recvfrom(100)
print("Sender",sender,"replied with time:",
      time.decode("utf-8"))
s.close()
```

# TCP Fibonacci Service - Iterative

```
# Server

from socket import *

def fib(n):
    if n == 0: return 0
    elif n == 1: return 1
    else: return fib(n-1)+fib(n-2)

def handle_request(client):
    buf = ""
    while True:
        d = client.recv(1).decode("utf-8")
        if d == "\n" or d == "\0":
            break
        buf += d
    number = int(buf)
    fibNumber = fib(number)
    client.send(str(fibNumber).encode("utf-8"))
    client.close()

if __name__ == "__main__":
    s = socket(AF_INET, SOCK_STREAM)
    s.setsockopt(SOL_SOCKET, SO_REUSEADDR, 1)
    s.bind(("0.0.0.0", 9000))
    s.listen(1)
    try:
        while True:
            client_socket, peer = s.accept()
            handle_request(client_socket)
    except KeyboardInterrupt:
        pass
    finally:
        s.close()
```

```
# Client

from socket import *
import datetime

start = datetime.datetime.now()
n = 37
s = socket(AF_INET, SOCK_STREAM)
s.connect(("localhost", 9000))
s.send(f"{n}\n".encode("utf-8"))
r = int(s.recv(10))
stop = datetime.datetime.now()
d = (stop - start).total_seconds()
print(f"fib({n})={r} ({d}s calc time)")
s.close()
```

# TCP Fibonacci Service - Forked

```
# Server

from socket import *
import os

s = socket(AF_INET, SOCK_STREAM)
s.setsockopt(SOL_SOCKET, SO_REUSEADDR, 1)
s.bind(("0.0.0.0", 9000))
s.listen(5)

# Avoid zombies
signal.signal(signal.SIGCHLD, signal.SIG_IGN)

try:
    while True:
        client_socket, peer = s.accept()
        child_pid = os.fork()

        # Handle fork
        if child_pid == 0:
            handle_request(client_socket)
            break
        else:
            client_socket.close()
except KeyboardInterrupt:
    pass
finally:
    s.close()
```

## Changes

- Fork after accept.
- `fork`: Creates a child process as copy of current process.
- Now: Two processes at the same line of code.
- Returns: 0 (inside child process), non-zero (parent process).

## Details

- Handle `SIGCHLD`, otherwise we get zombie processes.
- Close `child_socket` in parent, otherwise it will never be released.



# TCP Fibonacci Service - Threaded $\rightleftarrows$

```
from socket import *
from threading import Thread

s = socket(AF_INET, SOCK_STREAM)
s.setsockopt(SOL_SOCKET, SO_REUSEADDR, 1)
s.bind(("0.0.0.0", 9000))
s.listen(5)
try:
    while True:
        client_socket, peer = s.accept()
        t = Thread(target=handle_request,
                    args=(client_socket,))
        t.start()
except KeyboardInterrupt:
    pass
finally:
    s.close()
```

## Changes


- Instead of `fork`, create a `thread` and give it the request method to run and `client_socket` to process.
- More lightweight and faster than `fork`, because only a second execution thread is generated. Memory is shared (leading to the usual problems)!

# Concurrency and Parallelism - Advice

⚠ Not a central topic of this course, but **very** important for network engineers.

- Consider "[Nebenläufige Programmierung](#)" (basic Computer Science lecture, offered in summer semesters by Prof. Dr.-Ing. Holger Hermanns).
- Read a good book on concurrency.

## Java Concurrency in Practice



## Seven Concurrency Models in Seven Weeks: When Threads Unravel

The Pragmatic Programmers

Seven Concurrency Models  
in Seven Weeks  
When Threads Unravel



Paul Butcher

Series editor: Bruce A. Tate  
Development editor: Jacquelyn Carter

# Block or Not

# (Non-)Blocking I/O

## ⚡ Problem

- Some networking operations are blocking:
  - Examples: `send`, `recv`, `accept`, `connect`.
  - Program waits until operation succeeded.
  - Might wait forever if sender/receiver is not responsive.
- Example:
  - Client connects to server.
  - Server starts receiving, but client never sends.
  - Server waits forever.

## 💡 Solutions

A) Use blocking I/O. Use Multithreading. Spawn threads or child processes (see before).

B) Use blocking I/O. Specify Timeout. Try to send/recv for X seconds, abort afterwards.

C) Use non-blocking I/O. Call to blocking operation immediately returns an error if operation would have blocked; retry later.

D) I/O multiplexing. Check which socket "is ready" to do certain actions.

# Not to block (for too long)

## 🕒 Timeouts

- Timeouts specified on per-socket basis.
- `s.settimeout(timeout)`.
- `timeout`: timeout in seconds (float).

```
import socket

s = socket.socket(socket.AF_INET,
                  socket.SOCK_STREAM, 0)

s.settimeout(0.5)

try:
    s.connect(("127.0.0.1", 4567)) # try connect
    s.recv(1024)                  # try receive
except socket.timeout:
    print("Failed to connect/receive: TIMEOUT!")
s.close()
```

## ▶▶ Non-blocking operation

- Tell the OS to immediately return, even with no data (and hence error).
  - `s.setblocking(flag)`
  - `flag`: enable (`flag=1`) or disable (`flag=0`) blocking
- If successful, return as usual.
- But if socket is "not ready", throw exception:
  - Program has to retry later (`EAGAIN`, `EWOULDBLOCK`).
  - Unclean design: keep polling for non-blocking operation.

# Non-Blocking: Example

```
import socket
import errno
import time

s = socket.socket(socket.AF_INET,
                  socket.SOCK_STREAM,
                  0)

print("connecting!")
s.connect(("smtp.uni-saarland.de", 25))
print("connected!")
s.setblocking(0)

while True:
    try:
        print("received: {}".format(
            s.recv(1024).decode('utf-8')))
        break
    except socket.error as e:
        err = e.args[0]
        if err == errno.EAGAIN:
            print("EAGAIN")
        if err == errno.EWOULDBLOCK:
            print("EWOULDBLOCK")
        time.sleep(0.01) # avoid busy loop
```

```
$ python nonblocking.py
connecting!
connected!
EAGAIN
EWOULDBLOCK
EAGAIN
EWOULDBLOCK
EAGAIN
EWOULDBLOCK
EAGAIN
EWOULDBLOCK
EAGAIN
EWOULDBLOCK
received: 220 triton.rz.uni-saarland.de ESMTP
Sendmail; Fri, 2 Jun 2017 17:17:35 +0200
```

# I/O Multiplexing I

- Imagine you run a server:
  - Server handles hundreds of client connections.
  - Any client could send data at any time.
  - But you do not know which one, so you have to poll.
- Inefficient example:

```
for s in sockets: # iterate over clients
    try:
        # try to receive from this client
        data = s.recv()
    except socket.timeout:
        # nothing to receive: next client
        continue
```

- Query OS which sockets are "ready". OS returns sockets that...
  - ... can be read from.
  - ... can be sent to.
  - ... have an exception.
- Command: `select(rdrs, wrtrs, errs, tmout)`
  - `rdrs`: list of sockets to check if they are readable.
  - `wrtrs`: list of sockets to check if they are writeable.
  - `errs`: list of sockets to check if they have an I/O error.
  - `tmout`: timeout in seconds for `select()` to return.
  - Note: Select also accepts other I/O objects than sockets.

# I/O Multiplexing II

Sockets are readable if:

- There is at least one byte to read.
- The other end has terminated the connection (EOF).
- The socket is listening and `accept()` would not block.
- A socket error is pending.

Sockets are writeable if:

- Kernel socket buffers still fit at least one byte.
- The write half of the connection is closed(!).
- Non-blocking connect has completed.
- A socket error is pending.

```
while True:
    rds = wrs = errs = None

    try:
        (rds, wrs, errs) = select.select(sockets,
                                          [],
                                          sockets,
                                          0.2)

    except socket.timeout:
        continue # no socket ready or disrupted

    for s in rds: # read available sockets
        data = s.recv() # recv will *not* block

    for s in errs: # close sockets with errors
        s.close()
```



# More on Sockets

# Options

Manipulate socket behaviour (via `socket.setsockopt()`).

Socket Option	Description
SO_KEEPALIVE	Probe "silent" TCP connections if peer still reactive
SO_LINGER	Change socket close behaviour (flush/wait)
SO_RCVBUF	Size of receive buffer (default 4kB or more)
SO_SNDBUF	Size of send buffer (default 4kB or more)
SO_RCVLOWAT	Minimum number of bytes to recv (default = 1)
SO_SNDLOWAT	Minimum number of bytes to send (default = 1)
SO_REUSEADDR	Bind to port even if established connections exist to that port (e.g. quick restart of a server)
IP_TTL	Specify initial TTL of an IP packet
IP_RECVIF	Return on which interface a segment was received

# Retrieve Hostnames and Service Ports

## Hostnames:

- Use `socket.gethostbyname(hostname)` to translate the hostname to an IPv4 address.

Example: `gethostbyname('www.nt.uni-saarland.de')` -> 134.96.7.186

- DNS and contents of `/etc/hosts` are consumed by this:

```
# /etc/hosts
# The following lines are desirable for IPv4 capable hosts
127.0.0.1        localhost

# 127.0.1.1 is often used for the FQDN of the machine
127.0.1.1        thishost.mydomain.org  thishost
192.168.1.10     foo.mydomain.org       foo
192.168.1.13     bar.mydomain.org       bar
```

## Services:

- Use `socket.getservbyname(servicename[, protocolname])` to retrieve a protocol to a well-known port (for tcp or udp as `protocolname`).

Example: `getservbyname('ssh')` -> 22

- Mapping can be found (and edited) in `/etc/services`.

# Retrieve Address Information

**⚠ Problem:** Previously mentioned functions only work with IPv4 and require multiple calls to solve one task.

- New call combines all in one command:  
`socket.getaddrinfo(host, port, family=0, type=0, proto=0, flags=0)`
  - `host`: Domain Name, String representation of IPv4/v6 or None.
  - `port`: String service name (e.g. http).
  - `family`, `proto`, `flags` can be used to narrow results.
- Returns: `(family, type, proto, canonname, sockaddr)`
- `sockaddr` can be directly passed to `socket.connect()`.

# IP Protocols

- The file `/etc/protocols` lists the protocols (and their numbers) that are payloads of the IP layer.
  - e.g. TCP(6), UDP(17), ICMP(1), ...
- The number specified there is put into the `protocol` field in the IPv4 header and the `next_header` field in IPv6.
- Python's `socket` namespace includes `IPPROTO_*` constants that map to the same values.

# Wrap-Up

## Take-Home Messages

- Server design impacts performance. Chose your primitives well.
- Blocking calls have to be handled to avoid stalls.
- Functions `gethostbyname()` etc. allow you to utilize other protocols (e.g. DNS).

## Action Points

- **Download** and **solve** the task sheet.