

# Linux Networking

## Unit 18 - Hands-On Networking - 2018

*Prof. Dr.-Ing. Thorsten Herfet, [Andreas Schmidt](#), Pablo Gil Pereira*

**Telecommunications Lab, Saarland Informatics Campus, 26th Feb. 2018**

# Iproute2

# iproute2

- Manipulates network subsystem of the Linux kernel
- Replaces many traditional commands, e.g.:
  - ifconfig
  - route
  - arp
  - ...
- Unlike the tools above gives you nearly full control over linux network settings
- Functionality split up into sub-commands
- A little more complicated, but very well worth the learning curve

```
# Get help
$ ip help
...

$ ip <subcommand> help
...
```

# "Link" Layer

The `link` subcommand can be used to manipulate **interface / link settings**.

```
# show all interfaces / links of the current host
$ ip link show

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00

2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT \
   group default qlen 1000
   link/ether 70:54:d2:7b:7a:db brd ff:ff:ff:ff:ff:ff

3: eth1: <BROADCAST,MULTICAST> mtu 1500 qdisc pfifo_fast state DOWN mode DEFAULT group default \
   qlen 1000
   link/ether 70:54:d2:7b:7a:dc brd ff:ff:ff:ff:ff:ff
```

| Flag       | Description   |
|------------|---|
| UP         | The interface is enabled                                |
| LOWER_UP   | The lower layer is ready (e.g. "a cable is plugged in") |
| NO-CARRIER | The physical layer is not ready.                        |
| BROADCAST  | Interface supports broadcasts                           |
| MULTICAST  | Interface supports multicast                            |

# "Link" Layer

## Manipulating Link settings

```
# Enable / Disable an interface
$ ip link set up dev eth0
$ ip link set down dev eth0

# Changing the MTU
$ ip link set mtu 1492 dev eth0

# Enable / Disable promiscuous mode
$ ip link set promisc on dev eth0
$ ip link set promisc off dev eth0

# Set a new MAC address
$ ip link set dev eth0 address 02:ab:00:1f:2b:34

# Add a more human-friendly description to an interface
$ ip link set dev eth0 alias "Local Network"

$ ip link show dev eth0
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT \
    group default qlen 1000
    link/ether 70:54:d2:7b:7a:db brd ff:ff:ff:ff:ff:ff
    alias Local Network
```

# Linux and its interfaces

Nearly all network functions in Linux are represented by interfaces

## Physical Interfaces

- Represent a physical point of attachment to a network.  
i.e. a network card
- Data sent to the interface is put on the wire and the other way round.

## Virtual Interfaces

- Logical abstraction of a specific function.
- Data sent to a virtual interface may not go to the physical network.
- Data sent to a virtual interface may be processed in some way before going to the network.

# Virtual Interfaces | Example

## VLANs

A networked host can belong to **more than one VLAN** requiring it to **tag outgoing data frames**. In Linux different VLANs are represented by virtual interfaces. All data sent to this interface will be **tagged and sent out the underlying link**.

```
$ ip link
...
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT \
    group default qlen 1000
    link/ether 70:54:d2:7b:7a:db brd ff:ff:ff:ff:ff:ff

# Add VLAN 10 to interface eth0
$ ip link add name vlan10 link eth0 type vlan id 10

$ ip link
...
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT \
    group default qlen 1000
    link/ether 70:54:d2:7b:7a:db brd ff:ff:ff:ff:ff:ff
3: vlan10@eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT
    group default
    link/ether 70:54:d2:7b:7a:db brd ff:ff:ff:ff:ff:ff

# Remove the VLAN
$ ip link del dev vlan10
```

# Network Layer

The `address` subcommand can modify network layer settings, i.e. addresses.

```
# Show current IP configuration
$ ip address show

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP \
    group default qlen 1000
    link/ether 70:54:d2:7b:7a:db brd ff:ff:ff:ff:ff:ff
    inet 134.96.86.110/25 brd 134.96.86.127 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::7254:d2ff:fe7b:7adb/64 scope link
        valid_lft forever preferred_lft forever
```

```
# Show IPv6 Addresses only
$ ip -6 address show

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qlen 1000
    inet6 fe80::7254:d2ff:fe7b:7adb/64 scope link
        valid_lft forever preferred_lft forever
```



# Network Layer

## Manipulating IP Addresses

```
# Set an IP address
$ ip address add 10.0.0.1/24 dev eth1
$ ip address add 192.0.2.1/24 dev eth1
$ ip address add fd1e:1186:9d39::1/64 dev eth1

$ ip address
...
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 70:54:d2:7b:7a:dc brd ff:ff:ff:ff:ff:ff
    inet 10.0.0.1/24 brd 10.0.0.255 scope global eth1
        valid_lft forever preferred_lft forever
    inet 192.0.2.1/24 brd 192.0.2.255 scope global eth1
        valid_lft forever preferred_lft forever
    inet6 fd1e:1186:9d39::1/64 scope link
        valid_lft forever preferred_lft forever
    inet6 fe80::7254:d2ff:fe7b:7adc/64 scope link
        valid_lft forever preferred_lft forever

# Remove IP address from interface
$ ip address del 192.0.2.1/24 dev eth1

# Remove all addresses from interface
$ ip address flush dev eth1
```

# Network Layer

## Route Management

```
# Show the current routing table
$ ip route show

default via 192.0.2.1 dev eth0 proto static
192.0.2.0/25 dev eth0 proto kernel scope link src 192.0.2.11 metric 1

# Add a route via a gateway (i.e. router)
$ ip route add 192.0.2.128/25 via 192.0.2.1

# Change an existing route
$ ip route change 192.0.2.128/25 via 192.0.2.2

# Delete a route
$ ip route delete 192.0.2.128/25 via 192.0.2.2

# Add a default route
$ ip route add default via 192.0.2.1

# Discard traffic to unwanted destinations, e.g. known malicious hosts
$ ip route add blackhole 203.0.113.0/24 # drop silently
$ ip route add unreachable 203.0.113.0/24 # ICMP Host Unreachable
$ ip route add prohibit 203.0.113.0/24 # ICMP Administratively Prohibited
$ ip route add throw 203.0.113.0/24 # ICMP Network Unreachable
```

# Network Layer

Think about the neighbors ...

```
# Show the ARP / NDP neighbor cache
$ ip neighbor show

134.96.86.96 dev eth0 lladdr 30:05:5c:53:d3:86 STALE
134.96.86.1 dev eth0 lladdr 00:23:5d:ff:e4:00 REACHABLE
fe80::468a:5bff:fe27:880b dev eth0 lladdr 44:8a:5b:27:88:0b REACHABLE

# Clear the Cache
$ ip neighbor flush

# Manually add an entry to the Cache
$ ip neighbor add 192.0.2.14 lladdr 02:4a:bb:68:9c:d0
```

# Making it permanent

Iproute2 only changes the kernel's runtime settings.

```
$ nano /etc/network/interfaces
```

```
# Start interface automatically
auto eth0

# Use DHCP ...
iface eth0 inet dhcp

# ... or use a static IP config
iface eth0 inet static
    address 192.0.2.33
    netmask 255.255.255.0
    gateway 192.0.2.1           # default gateway

# VLAN 10 on interface eth0
iface vlan10 inet static
    address 10.0.0.1
    netmask 255.255.255.0
    vlan-raw-device eth0
```

```
# Configure IPv6 ...

# ... using SLAAC
# No configuration necessary.
# Linux uses SLAAC by default

# ... using DHCPv6
iface eth0 inet6 dhcp

# ... statically
iface eth0 inet6 static
    address fd1e:1186:9d39::33
    netmask 64
    gateway fd1e:1186:9d39::1
```

```
# Enable / Disable interface listed in interfaces file
$ ifup eth0
$ ifdown eth0
```

# Linux Kernel Internals

# Sysctl

- Sysctl is a **config interface** to the **kernel's runtime settings**.
- Allows you to **tune** many aspects of the **running kernel**.
  - CPU scheduling
  - Memory management
  - **Network settings**
  - ...
- Individual settings are represented as **files under /proc/sys**.
- **sysctl-command** can be used to change settings.

# Sysctl | Examples

```
# Check value of a setting
$ sysctl net.ipv4.ip_forward
net.ipv4.ip_forward = 0

# Linux does not forward packets by default
# Enable routing for IPv4
$ sysctl -w net.ipv4.ip_forward=1

# Enable routing for IPv6
$ sysctl -w net.ipv6.conf.all.forwarding=1

# Disable responses to ICMP echo requests
$ sysctl -w net.ipv4.icmp_echo_ignore_all=1

# Ignore IPv6 router advertisements received on eth0
$ sysctl -w net.ipv6.conf.eth0.ignore_ra=0
```

To change settings permanently you can add them to `/etc/sysctl.conf` or a file below `/etc/sysctl.d/` so they survive a system reboot.

Be careful when changing sysctl parameters. Some settings may break your network connection or make your system unstable.

Make sure you know what you are doing.

# Linux and Names



# Linux Hostname

```
# Display the hostname
$ hostname
honlinux

# Display the fully qualified hostname
$ hostname -f
honlinux.my.domain

# Set the hostname to "gateway"
# (until the next reboot)
$ hostname gateway
```

To set the hostname permanently, edit `/etc/hostname`.

```
echo "gateway" > /etc/hostname
```

# Local Resolver

- Part of the system **libc**
- Syscall **interface for name resolution**
- **Resolves names** to network addresses
- Specific features vary by implementation

# /etc/resolv.conf

Config file for the local resolver

- List nameservers
- List default search domains
- Resolver options

```
## /etc/resolv.conf

## List of nameservers
nameserver 134.96.7.5
nameserver 134.96.7.100

## Search suffixes
search nt.uni-saarland.de uni-saarland.de

## Options
# rotate: try a different server for every query
# timeout: number of seconds to wait for replies
options rotate timeout:5
```

# /etc/hosts

- Define hostname/IP pairs for the resolver
- Can be used to override names in the public DNS on a per-host basis
- One entry per line
- Each entry can contain **exactly one IP address** and **one or more hostnames**

```
## /etc/hosts

# Entries for the local host
127.0.0.1    honlinux.my.domain localhost
::1         honlinux.my.domain localhost

# A machine not in public DNS with an
# address we'd rather not have to remember
fd2f:5f10:b0a0:9277:aab:26a:0:18ba    dev-system    dev    d
```

# Dnsmasq

- Network infrastructure for **small networks**
  - DNS
  - DHCP / DHCPv6
  - IPv6 Router advertisements
  - Network boot
- Lightweight
- Easy to configure (`/etc/dnsmasq.conf`)

# Network Resolver (using Dnsmasq)

Just start Dnsmasq. It will use the local resolver to forward requests by default.

```
$ service dnsmasq start
```

## Checking the Resolver

- host (🐧)
- dig (🐧, 🍏)
- nslookup (🐧, 🍏, 🪟)

```
# Check the DNS resolver running on the local host
$ host www.uni-saarland.de 127.0.0.1
...

$ dig @127.0.0.1 www.uni-saarland.de
...

$ nslookup
> server 127.0.0.1
Default server: 127.0.0.1
Address: 127.0.0.1#53
> www.uni-saarland.de
...
```

# DNS Server (using Dnsmasq)

Dnsmasq answers DNS requests matching hostnames in the local `/etc/hosts`.

```
## /etc/hosts
...
192.0.2.10    server.my.domain
...
```

```
$ dig @127.0.0.1 server.my.domain
...
;; ANSWER SECTION:
server.my.domain.      170      IN      A       192.0.2.10
...
```

# Dynamic Addressing (using Dnsmasq)

```
## /etc/dnsmasq.conf

# Define an IPv4 address pool (192.0.2.10-20, 12h lease time)
dhcp-range=192.0.2.10, 192.0.2.20, 12h

# DHCPv6 (201 addresses with /64 prefix, 12h lease time)
dhcp-range=fd1e:1186:9d39::100, fd1e:1186:9d39::200, 64, 12h

# SLAAC (enable router advertisements, no DHCP)
# (Note for the tutorials: SLAAC breaks VPCS!)
dhcp-range=fd1e:1186:9d39::, ra-stateless

# Fixed IP address for a specific host
dhcp-host=11:22:33:44:55:66, 192.0.2.44

# Dnsmasq announces itself as the default gateway, DNS and NTP server by default
```

```
# Don't forget to restart Dnsmasq for changes to take effect
$ service dnsmasq restart
```

```
# Make sure Dnsmasq is started on system boot
$ rc-update add dnsmasq default
```



# iptables | Examples

```
# DNAT (Change destination address)
iptables -t nat -A PREROUTING -s 192.0.2.12 -p tcp --dport 80 -j DNAT --to-destination 203.0.113.44
```

```
# SNAT (Change Source address)
iptables -t nat -A POSTROUTING -o eth0 -j SNAT --to-source 198.51.100.1
```

```
# masquerading NAT (Change source address to that of outgoing interface)
iptables -t nat -A POSTROUTING -o eth0 -j MASQUERADE
```

```
# To filter IPv6 traffic, Linux uses a different set of tables.
# You can use ip6tables to manipulate them.
$ ip6tables -s fd1e:1186:9d39::1aa4 -j DROP
```

```
# Once you checked your current ruleset, you can persist it ...
$ service iptables save
$ service ip6tables save

# ... and automatically load it on system boot
$ rc-update add iptables default
$ rc-update add ip6tables default
```

# Wrap-Up

## Further Reading

- [Iproute2 Cheat Sheet](#)
- `man interfaces`
- [Sysctl Documentation](#)
- [Sysctl Network Settings Documentation](#)
- `man resolv.conf`
- [Dnsmasq Website](#)
- `man dnsmasq`